

URL: https://scikit-learn.org/stable/modules/cross_validation.html#cross-validation-and-model-selection As shown in the tutorial, create a notebook that replicates and implements at least four different cross-validation methods. Then pick two cross-validation methods to compare the performance of your best-performing SVM, Decision tree, AdaBoost, and Random Forest models on the breast cancer data from HW5.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.pipeline import make_pipeline
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, ShuffleSplit, RepeatedKFold, cross_val_score, GridSearchCV
from sklearn import datasets, svm

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier, RandomForestClassifier
```

```
In [2]: %matplotlib inline

# Turn off warnings completely for the Notebook
import warnings
warnings.filterwarnings('ignore')
```

Part 1. Replicates four types of cross-validation methods

```
In [3]: # Load iris data
X, y = datasets.load_iris(return_X_y=True)
X.shape, y.shape
```

```
Out[3]: ((150, 4), (150,))
```

```
In [4]: # split train-test subsets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0, stratify=y)
```

```
In [5]: # fit SVM model
clf = make_pipeline(StandardScaler(), svm.SVC())
clf.fit(X_train, y_train)
clf.score(X_test, y_test)
```

```
Out[5]: 0.9777777777777777
```

```
In [6]: # cross-validation computing
# Method 1. computing the score 5 consecutive times (with different splits each time)
scores_1 = cross_val_score(clf, X, y, cv=5)
scores_1
```

```
Out[6]: array([0.96666667, 0.96666667, 0.96666667, 0.93333333, 1.          ])
```

```
In [7]: print('%.2f accuracy with a standard deviation of %.2f'%(scores_1.mean(), scores_1.std()))
```

```
0.97 accuracy with a standard deviation of 0.02
```

```
In [8]: # Method 2. ShuffleSplit
n_samples = X.shape[0]
cv = ShuffleSplit(n_splits=5, test_size=0.3, random_state=0)
scores_2 = cross_val_score(clf, X, y, cv=cv)
scores_2
```

```
Out[8]: array([0.97777778, 0.93333333, 0.95555556, 0.93333333, 0.97777778])
```

```
In [9]: print('%.2f accuracy with a standard deviation of %.2f'%(scores_2.mean(), scores_2.std()))
```

```
0.96 accuracy with a standard deviation of 0.02
```

```
In [10]: # Method 3. use an iterable yielding (train, test) splits
def custom_cv_2folds(X):
    n = X.shape[0]
    i = 1
    while i <= 2:
        idx = np.arange(n * (i - 1) / 2, n * i / 2, dtype=int)
        yield idx, idx
        i += 1

custom_cv = custom_cv_2folds(X)
scores_3 = cross_val_score(clf, X, y, cv=custom_cv)
print('%.2f accuracy with a standard deviation of %.2f'%(scores_3.mean(), scores_3.std()))
```

```
0.99 accuracy with a standard deviation of 0.01
```

```
In [11]: # Method 4. Repeated K-Fold
rkf = RepeatedKFold(n_splits=2, n_repeats=2, random_state=123)
scores_4 = cross_val_score(clf, X, y, cv=rkf)
print('%.2f accuracy with a standard deviation of %.2f'%(scores_4.mean(), scores_4.std()))
```

0.94 accuracy with a standard deviation of 0.01

Part 2. pick two cross-validation methods to compare the performance of your best-performing SVM, Decision tree, AdaBoost, and Random Forest models on the breast cancer data from HW5

```
In [12]: # Load data
wbc = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/wdbc.data', header=None)
```

```
In [13]: # check for any NA
wbc.isna().sum().sum()
```

```
Out[13]: 0
```

```
In [14]: wbc.head()
```

```
Out[14]:
```

	0	1	2	3	4	5	6	7	8	9	...	22	23	24	25	26	27	28	29	30	31
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	...	25.38	17.33	184.60	2019.0	0.1622	0.6656	0.7119	0.2654	0.4601	0.11890
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...	24.99	23.41	158.80	1956.0	0.1238	0.1866	0.2416	0.1860	0.2750	0.08902
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...	23.57	25.53	152.50	1709.0	0.1444	0.4245	0.4504	0.2430	0.3613	0.08758
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	...	14.91	26.50	98.87	567.7	0.2098	0.8663	0.6869	0.2575	0.6638	0.17300
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	...	22.54	16.67	152.20	1575.0	0.1374	0.2050	0.4000	0.1625	0.2364	0.07678

5 rows × 32 columns

```
In [15]: # split as features and target
X, y = wbc.iloc[:, 2:], wbc.iloc[:, 1]
X.shape, y.shape
```

```
Out[15]: ((569, 30), (569,))
```

```
In [16]: # Divide the data into train (80%) and test (20%)
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1, stratify=y, test_size=0.2)
```

```
In [17]: clf1 = DecisionTreeClassifier(random_state=123)
clf2 = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), random_state=123)
clf3 = RandomForestClassifier(random_state=123)
clf4 = svm.SVC(random_state=123)

pipe4 = Pipeline([['sc', StandardScaler()], ['clf', clf4]])
```

```
In [18]: ## Evaluating and tuning the ensemble classifier with GridSearchCV =====
# get a basic idea of how we can access the individual parameters inside a GridSearch object:
pipe4.get_params()
```

```
Out[18]: {'memory': None,
'steps': [['sc', StandardScaler()], ['clf', SVC(random_state=123)]],
'verbose': False,
'sc': StandardScaler(),
'clf': SVC(random_state=123),
'sc__copy': True,
'sc__with_mean': True,
'sc__with_std': True,
'clf__C': 1.0,
'clf__break_ties': False,
'clf__cache_size': 200,
'clf__class_weight': None,
'clf__coef0': 0.0,
'clf__decision_function_shape': 'ovr',
'clf__degree': 3,
'clf__gamma': 'scale',
'clf__kernel': 'rbf',
'clf__max_iter': -1,
'clf__probability': False,
'clf__random_state': 123,
'clf__shrinking': True,
'clf__tol': 0.001,
'clf__verbose': False}
```

```
In [19]: # tune the parameters via a grid search
params_1 = {'criterion':['gini','entropy'],
            'max_depth':[5,10,20]}
params_2 = {'n_estimators':[50,500,1000],
            'learning_rate':[0.5, 1.0, 5.0],
            'algorithm':['SAMME.R','SAMME']}
params_3 = {'criterion':['gini','entropy'],
            'max_depth':[5,10,20],
            'min_samples_split':[20,50,100],
            'min_samples_leaf':[5,10],
            'oob_score':[True, False]}
params_4 = {'clf__C':[0.001, 0.1, 100.0],
            'clf__kernel':['linear', 'poly', 'rbf', 'sigmoid', 'precomputed'],
            'clf__gamma':['scale', 'auto'],
            'clf__shrinking':[True,False]}
```

```

In [20]: # Define function for best parameters search
def find_best_param(clf,param):
    grid = GridSearchCV(estimator=clf, param_grid=param, cv=10)
    grid.fit(X_train, y_train)

    print(grid.best_score_)
    print(grid.best_params_)
    print(grid.best_estimator_)

In [21]: # best performing decision tree
find_best_param(clf1,params_1)

0.9561835748792271
{'criterion': 'gini', 'max_depth': 10}
DecisionTreeClassifier(max_depth=10, random_state=123)

In [22]: clf1_b = DecisionTreeClassifier(criterion='gini', max_depth=10, random_state=123)

In [23]: # best performing adaBoost classifier
find_best_param(clf2,params_2)

0.9517391304347826
{'algorithm': 'SAMME.R', 'learning_rate': 0.5, 'n_estimators': 50}
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), learning_rate=0.5,
                    random_state=123)

In [24]: clf2_b = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(),
                                   algorithm='SAMME.R', learning_rate=0.5, n_estimators=50,
                                   random_state=123)

In [25]: # best performing random forest classifier
find_best_param(clf3, params_3)

0.953816425120773
{'criterion': 'entropy', 'max_depth': 5, 'min_samples_leaf': 5, 'min_samples_split': 20, 'oob_score': True}
RandomForestClassifier(criterion='entropy', max_depth=5, min_samples_leaf=5,
                       min_samples_split=20, oob_score=True, random_state=123)

In [26]: clf3_b = RandomForestClassifier(criterion='entropy', max_depth=5, min_samples_leaf=5,
                                       min_samples_split=20, oob_score=True, random_state=123,
                                       max_features='auto')

In [27]: # best performing SVM
find_best_param(pipe4, params_4)

0.9802415458937197
{'clf__C': 0.1, 'clf__gamma': 'scale', 'clf__kernel': 'linear', 'clf__shrinking': True}
Pipeline(steps=[('sc', StandardScaler()),
                 ('clf', SVC(C=0.1, kernel='linear', random_state=123))])

In [28]: pipe4_b = Pipeline(steps=[('sc', StandardScaler()),
                                   ('clf', svm.SVC(C=0.1, kernel='linear', gamma='scale',
                                   shrinking=True, random_state=123))])

In [29]: clf_labels = ['Decision tree', 'adaBoost', 'RandForest', 'SVM']

In [30]: # cross-validation I == 10-fold cross validation

print('10-fold cross validation:\n')
for clf, label in zip([clf1_b, clf2_b, clf3_b, pipe4_b], clf_labels):
    scores = cross_val_score(estimator=clf,
                             X=X_train,
                             y=y_train,
                             cv=10)

    print("Accuracy with Standard Deviation: %0.2f (+/- %0.2f) [%s]"
          % (scores.mean(), scores.std(), label))

10-fold cross validation:

Accuracy with Standard Deviation: 0.96 (+/- 0.02) [Decision tree]
Accuracy with Standard Deviation: 0.95 (+/- 0.03) [adaBoost]
Accuracy with Standard Deviation: 0.95 (+/- 0.02) [RandForest]
Accuracy with Standard Deviation: 0.98 (+/- 0.02) [SVM]

In [31]: # cross-validation II == ShuffleSplit cross validation
n_samples = X_train.shape[0]
cv = ShuffleSplit(n_splits=10, test_size=0.3, random_state=0)

In [32]: print('shuffle-split cross validation:\n')
for clf, label in zip([clf1_b, clf2_b, clf3_b, pipe4_b], clf_labels):
    scores = cross_val_score(estimator=clf,
                             X=X_train,
                             y=y_train,
                             cv=cv)

    print("Accuracy with Standard Deviation: %0.2f (+/- %0.2f) [%s]"
          % (scores.mean(), scores.std(), label))

```

shuffle-split cross validation:

Accuracy with Standard Deviation: 0.93 (+/- 0.02) [Decision tree]

Accuracy with Standard Deviation: 0.93 (+/- 0.02) [adaBoost]

Accuracy with Standard Deviation: 0.95 (+/- 0.02) [RandForest]

Accuracy with Standard Deviation: 0.98 (+/- 0.01) [SVM]

In []: